

```
kate_kosareva@kos-Inspiron-5570:~/Documents/MADE/MLOps$ kubectl describe pods
Name:          online-inference-resources
Namespace:     default
Priority:       0
Service Account: default
Node:          docker-desktop/192.168.65.4
Start Time:    Tue, 13 Dec 2022 10:32:01 +0100
Labels:        app=online-inference
Annotations:    <none>
Status:        Running
IP:            10.1.0.14
IPs:           10.1.0.14
Containers:
  online-inference:
    Container ID:  docker://89c8189e57f0caea49e70ca924a9be13b69dea72ad5d0b1c021da0a6066f2755
    Image:         koluzajka/mlops_made:v.01
    Image ID:      docker-pullable://evgerher/ml-in-prod-inference@sha256:e3ef0b14e8ca296169a025ec7a615622f1d675445ab9a8e429c296367ba08782
    Port:          5050/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Tue, 13 Dec 2022 10:32:53 +0100
    Ready:         True
    Restart Count:  0
    Limits:
      cpu:          500m
      memory:       2Gi
    Requests:
      cpu:          250m
      memory:       1Gi
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-cxtw6 (ro)
Conditions:
  Type             Status
  Initialized       True
  Ready             True
  ContainersReady   True
  PodScheduled      True
Volumes:
  kube-api-access-cxtw6:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:  kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:    true
QoS Class:         Burstable
Node-Selectors:     <none>
Tolerations:        node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                    node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   28m   default-scheduler Successfully assigned default/online-inference-resources to docker-desktop
  Normal  Pulling     28m   kubelet       Pulling image "evgerher/ml-in-prod-inference:0.0.1"
```

```

  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   28m   default-scheduler Successfully assigned default/online-inference-resources to docker-desktop
  Normal  Pulling     28m   kubelet       Pulling image "evgerher/ml-in-prod-inference:0.0.1"
  Normal  Pulled      27m   kubelet       Successfully pulled image "evgerher/ml-in-prod-inference:0.0.1" in 49.930184475s
  Normal  Created     27m   kubelet       Created container online-inference
  Normal  Started     27m   kubelet       Started container online-inference
  Normal  Killing     8s    kubelet       Container online-inference definition changed, will be restarted

Name:          online-inference2
Namespace:     default
Priority:       0
Service Account: default
Node:          <none>
Labels:        app=online-inference2
Annotations:    <none>
Status:        Pending
IP:            <none>
IPs:           <none>
Containers:
  online-inference:
    Image:         koluzajka/mlops_made:v2
    Port:          5050/TCP
    Host Port:     0/TCP
    Limits:
      cpu:          500m
      memory:       2Gi
    Requests:
      cpu:          250m
      memory:       1Gi
    Liveness:       http-get http://:5050/health delay=40s timeout=1s period=3s #success=1 #failure=3
    Readiness:      http-get http://:5050/ready delay=10s timeout=1s period=3s #success=1 #failure=3
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-tlkdl (ro)
Conditions:
  Type             Status
  PodScheduled      False
Volumes:
  kube-api-access-tlkdl:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:  kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:    true
QoS Class:         Burstable
Node-Selectors:     <none>
Tolerations:        node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                    node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Warning  FailedScheduling 3m20s default-scheduler 0/1 nodes are available: 1 Insufficient memory. preemption: 0/1 nodes are available: 1 No preemption victims
```